

Name:N.Sai Sreenu

Reg No:192365001

C04-AT1

MLA0201-Fundamentals Of Machine Learning

Code Challenge

Question 1: Bayesian Network for Disease Prediction

Problem Statement

A healthcare organization wants to predict whether a patient has Diabetes based on the symptoms Obesity and High Blood Sugar. The goal is to construct a Bayesian Network, define suitable Conditional Probability Tables (CPTs), and perform probabilistic inference to find the probability of Diabetes when both symptoms are present.

Approach

- Model the causal structure $\text{Obesity} \rightarrow \text{Diabetes} \leftarrow \text{High Blood Sugar}$ using pgmpy's Bayesian Network class.
- Define prior CPTs for Obesity and High Blood Sugar, and a conditional CPT for Diabetes given both parents.
- Validate the network with `check_model()` to ensure the CPTs are consistent with the graph structure.
- Use the Variable Elimination algorithm to perform exact inference.
- Query $P(\text{Diabetes} \mid \text{Obesity} = \text{Yes}, \text{HighBloodSugar} = \text{Yes})$ and print the resulting probability distribution.

Python Code

```
# Question 1: Bayesian Network for Disease Prediction
# Predict Diabetes based on Obesity and High Blood Sugar
# Library required: pip install pgmpy

from pgmpy.models import BayesianNetwork
from pgmpy.factors.discrete import TabularCPD
from pgmpy.inference import VariableElimination

# Step 1: Construct the Bayesian Network structure
# Obesity -> Diabetes <- HighBloodSugar
```

```

diabetes_model = BayesianNetwork([
    ('Obesity', 'Diabetes'),
    ('HighBloodSugar', 'Diabetes')
])

# Step 2: Define Conditional Probability Tables (CPTs)

# Prior probability of Obesity (0 = No, 1 = Yes)
cpd_obesity = TabularCPD(
    variable='Obesity', variable_card=2,
    values=[[0.7], [0.3]]      # P(Obesity=No)=0.7, P(Obesity=Yes)=0.3
)

# Prior probability of High Blood Sugar (0 = No, 1 = Yes)
cpd_sugar = TabularCPD(
    variable='HighBloodSugar', variable_card=2,
    values=[[0.65], [0.35]]    # P(HBS=No)=0.65, P(HBS=Yes)=0.35
)

# Conditional probability of Diabetes given Obesity and High Blood Sugar
# Order of parent combinations: (Obesity, HighBloodSugar)
# (0,0) (0,1) (1,0) (1,1)
cpd_diabetes = TabularCPD(
    variable='Diabetes', variable_card=2,
    values=[
        [0.95, 0.55, 0.60, 0.10], # P(Diabetes=No | Obesity, HBS)
        [0.05, 0.45, 0.40, 0.90]  # P(Diabetes=Yes | Obesity, HBS)
    ],
    evidence=['Obesity', 'HighBloodSugar'],
    evidence_card=[2, 2]
)

# Step 3: Add CPTs to the model
diabetes_model.add_cpds(cpd_obesity, cpd_sugar, cpd_diabetes)

# Step 4: Validate the model
assert diabetes_model.check_model(), "Bayesian Network is invalid!"

# Step 5: Perform inference using Variable Elimination
inference = VariableElimination(diabetes_model)

# Predict probability of Diabetes when BOTH symptoms are present

```


```

result = inference.query(
    variables=['Diabetes'],
    evidence={'Obesity': 1, 'HighBloodSugar': 1}
)

# Step 6: Display the predicted probability
print("Bayesian Network: Diabetes Prediction")
print("-----")
print(result)
prob_diabetes_yes = result.values[1]
print(f"\nPredicted Probability of Diabetes (Obesity=Yes, HighBloodSugar=Yes): "
      f"{prob_diabetes_yes:.4f}")

```

Sample Output



```

Bayesian Network: Diabetes Prediction
-----
+-----+-----+
| Diabetes | phi(Diabetes) |
+=====+=====+
| Diabetes(0) | 0.1000 |
+-----+-----+
| Diabetes(1) | 0.9000 |
+-----+-----+

Predicted Probability of Diabetes (Obesity=Yes, HighBloodSugar=Yes): 0.9000

```

Explanation of Output

The inference engine returns the posterior distribution over Diabetes given that both Obesity and High Blood Sugar are observed as 'Yes'. Based on the CPT defined in the model, the predicted probability of Diabetes being present is 0.90 (90%), reflecting the strong influence both risk factors have when they co-occur.

Question 2: Hidden Markov Model for Weather Prediction

Problem Statement

A weather forecasting agency records daily observations, while the true underlying weather condition is hidden. The task is to implement a Hidden Markov Model with defined transition and emission probabilities, and to predict the most likely sequence of hidden weather states for a given sequence of observations.

Approach

- Define the hidden states (Sunny, Rainy) and the observable symptoms (Walk, Shop, Clean).
- Specify the initial state distribution, the state transition matrix, and the emission (observation) probabilities.
- Implement the Viterbi algorithm from scratch to find the most probable path of hidden states.
- Run the algorithm on a sample observation sequence and back-trace the optimal state path.
- Print the predicted hidden state sequence together with its probability.

Python Code

```
# Question 2: Hidden Markov Model for Weather Prediction
# Hidden states: Sunny, Rainy  Observations: Walk, Shop, Clean
# (Observations Sunny/Cloudy/Rainy are what a sensor reports;
# here we use a classic Walk/Shop/Clean observation set to keep
# hidden states and observations clearly distinct.)
# Implemented from scratch using the Viterbi algorithm (no external HMM library needed)

import numpy as np

# Step 1: Define hidden states and observations
states = ['Sunny', 'Rainy']
observations = ['Walk', 'Shop', 'Clean']

# Step 2: Define initial state probabilities
start_prob = {'Sunny': 0.6, 'Rainy': 0.4}

# Step 3: Define transition probabilities (hidden state -> hidden state)
transition_prob = {
    'Sunny': {'Sunny': 0.7, 'Rainy': 0.3},
    'Rainy': {'Sunny': 0.4, 'Rainy': 0.6}
}

# Step 4: Define emission probabilities (hidden state -> observation)
```

```

emission_prob = {
    'Sunny': {'Walk': 0.6, 'Shop': 0.3, 'Clean': 0.1},
    'Rainy': {'Walk': 0.1, 'Shop': 0.4, 'Clean': 0.5}
}

# Step 5: Sequence of observed events
obs_sequence = ['Walk', 'Shop', 'Clean', 'Walk']

def viterbi(obs_seq, states, start_p, trans_p, emit_p):
    """Predict the most likely sequence of hidden states using the Viterbi algorithm."""
    V = [{}]
    path = {}

    # Initialization
    for s in states:
        V[0][s] = start_p[s] * emit_p[s][obs_seq[0]]
        path[s] = [s]

    # Recursion
    for t in range(1, len(obs_seq)):
        V.append({})
        new_path = {}

        for curr_state in states:
            best_prob, best_prev_state = max(
                (V[t - 1][prev_state] * trans_p[prev_state][curr_state] *
                 emit_p[curr_state][obs_seq[t]], prev_state)
                for prev_state in states
            )
            V[t][curr_state] = best_prob
            new_path[curr_state] = path[best_prev_state] + [curr_state]

        path = new_path

    # Termination: pick the state with highest final probability
    final_prob, final_state = max((V[len(obs_seq) - 1][s], s) for s in states)
    return final_prob, path[final_state]

# Step 6: Run Viterbi algorithm to predict hidden weather states
probability, best_path = viterbi(obs_sequence, states, start_prob,
```

```
transition_prob, emission_prob)
```

```
# Step 7: Print results
print("Hidden Markov Model: Weather Prediction")
print("-----")
print("Observation Sequence :", obs_sequence)
print("Predicted Hidden States :", best_path)
print(f"Probability of Predicted Sequence : {probability:.6f}")
```

Sample Output



```
Hidden Markov Model: Weather Prediction
-----
Observation Sequence : ['Walk', 'Shop', 'Clean', 'Walk']
Predicted Hidden States : ['Sunny', 'Rainy', 'Rainy', 'Sunny']
Probability of Predicted Sequence : 0.003110
```

Explanation of Output

For the observation sequence ['Walk', 'Shop', 'Clean', 'Walk'], the Viterbi algorithm computes the most probable hidden weather sequence as ['Sunny', 'Rainy', 'Rainy', 'Sunny'], with an overall path probability of approximately 0.0031. This means the model believes the day started Sunny, turned Rainy for two days, and became Sunny again.

Question 3: Markov Random Field for Image Denoising

Problem Statement

A computer vision company wants to remove noise from grayscale images using a Markov Random Field. Each pixel is treated as a node in a graph, and neighboring pixels are connected by undirected edges representing spatial dependency, which is used to smooth noisy pixel values.

Approach

- Represent a small 5x5 grayscale image as a NumPy array of pixel intensities.
- Build a 4-connected undirected grid graph with NetworkX, where each node corresponds to one pixel.
- Attach each pixel's intensity value as a node attribute in the graph.
- Simulate one MRF smoothing iteration: update every pixel as a weighted average of its own value and the mean of its neighboring pixel values.
- Display the original image and the updated (denoised) image side by side.

Python Code

```
# Question 3: Markov Random Field for Image Denoising
# Represent a small grayscale image as a graph (grid) and smooth it
# using the average of each pixel's neighbors (4-connected MRF).

import numpy as np
import networkx as nx

# Step 1: Represent a small noisy grayscale image (5x5) as a NumPy array
original_image = np.array([
    [52, 55, 61, 59, 79],
    [62, 59, 55, 104, 94],
    [63, 65, 66, 113, 144],
    [61, 65, 67, 120, 138],
    [59, 63, 68, 120, 133]
])

rows, cols = original_image.shape

# Step 2: Build a graph where each pixel is a node
# and edges connect 4-connected neighboring pixels
graph = nx.grid_2d_graph(rows, cols)

# Attach the pixel intensity as a node attribute
```

```

for (r, c) in graph.nodes():
    graph.nodes[(r, c)][‘value’] = original_image[r, c]

def mrf_smoothing_iteration(image, graph, weight_self=0.5, weight_neighbors=0.5):
    """
    Simulate one iteration of MRF-based smoothing.
    Each pixel's new value = weighted average of its own value
    and the mean value of its graph neighbors (undirected MRF edges).
    """
    new_image = image.copy().astype(float)

    for node in graph.nodes():
        r, c = node
        neighbor_values = [graph.nodes[n][‘value’] for n in graph.neighbors(node)]
        neighbor_mean = np.mean(neighbor_values)
        current_value = graph.nodes[node][‘value’]

        # MRF update rule: balance between observed pixel and neighborhood
        new_value = (weight_self * current_value) + (weight_neighbors * neighbor_mean)
        new_image[r, c] = new_value

    return new_image

# Step 3: Run one iteration of smoothing
updated_image = mrf_smoothing_iteration(original_image, graph)

# Step 4: Display original and updated pixel values
np.set_printoptions(precision=2, suppress=True)
print("Markov Random Field: Image Denoising")
print("-----")
print("Original Image (pixel values):")
print(original_image)
print("\nUpdated Image after 1 MRF smoothing iteration:")
print(updated_image)

```


Sample Output



```
Markov Random Field: Image Denoising
```

```
-----
```

```
Original Image (pixel values):
```

```
[[ 52  55  61  59  79]
 [ 62  59  55 104  94]
 [ 63  65  66 113 144]
 [ 61  65  67 120 138]
 [ 59  63  68 120 133]]
```

```
Updated Image after 1 MRF smoothing iteration:
```

```
[[ 55.25  56.17  58.67  70.17  77.75]
 [ 60.    59.12  63.75  92.12 101.5 ]
 [ 62.83  64.12  70.5   110.75 129.5 ]
 [ 61.67  64.5   73.38  114.75 135.17]
 [ 60.5   63.5   75.67  113.5   131.  ]]
```

Explanation of Output

The updated image shows each pixel value pulled slightly toward its local neighborhood average, which reduces sharp, noise-like variations while largely preserving the overall structure/edges of the image — exactly the smoothing behaviour expected from a single MRF update iteration.

Question 4: Conditional Random Field for Named Entity Recognition

Problem Statement

A customer support system needs to automatically identify Customer Names, Order IDs, and Product Names from free-text support emails. The task is to build sequential, word-level training data, train a Conditional Random Field model, and use it to label the entities in a new, unseen sentence.

Approach

- Create labelled training sentences as sequences of (word, POS tag, BIO label) tuples, using the BIO tagging scheme (B-CUST, I-CUST, B-ORDER, B-PROD, I-PROD, O).
- Engineer word-level features for each token: lower-cased form, suffix, casing, digit presence, POS tag, and features of the previous/next word.
-
- Convert a new, unlabelled sentence into the same feature representation and predict its label sequence.
- Print each word next to its predicted entity label in a clean tabular format.

Python Code

```
# Question 4: Conditional Random Field for Named Entity Recognition
# Identify Customer Names, Order IDs, and Product Names in support emails
# Library required: pip install sklearn-crfsuite

import sklearn_crfsuite
from sklearn_crfsuite import metrics

# Step 1: Create sequential training data
# Each sentence is a list of (word, POS-like-tag, label) tuples
train_sentences = [
    [("Hello", "O", "O"), ("John", "PROPN", "B-CUST"), ("Smith", "PROPN", "I-CUST"),
     ("ordered", "VERB", "O"), ("Order123", "NOUN", "B-ORDER"),
     ("for", "ADP", "O"), ("iPhone14", "NOUN", "B-PROD")],

    [("Dear", "O", "O"), ("support", "NOUN", "O"), ("my", "PRON", "O"),
     ("Order987", "NOUN", "B-ORDER"), ("for", "ADP", "O"),
     ("Samsung", "PROPN", "B-PROD"), ("TV", "NOUN", "I-PROD"),
     ("is", "AUX", "O"), ("delayed", "ADJ", "O")],

    [("Hi", "O", "O"), ("this", "PRON", "O"), ("is", "AUX", "O"),
     ("Emma", "PROPN", "B-CUST"), ("Watson", "PROPN", "I-CUST")],
```

```
    ("regarding", "ADP", "O"), ("OrderA45", "NOUN", "B-ORDER")],  
]
```

Step 2: Define word-level feature extraction functions

```
def word_features(sentence, i):  
    word, pos, _ = sentence[i]  
    features = {  
        'bias': 1.0,  
        'word.lower()': word.lower(),  
        'word[-3:]': word[-3:],  
        'word.isupper()': word.isupper(),  
        'word.istitle()': word.istitle(),  
        'word.isdigit()': word.isdigit(),  
        'has_digit': any(ch.isdigit() for ch in word),  
        'postag': pos,  
    }  
    if i > 0:  
        prev_word, prev_pos, _ = sentence[i - 1]  
        features.update({  
            '-1:word.lower()': prev_word.lower(),  
            '-1:word.istitle()': prev_word.istitle(),  
            '-1:postag': prev_pos,  
        })  
    else:  
        features['BOS'] = True # Beginning of sentence  
  
    if i < len(sentence) - 1:  
        next_word, next_pos, _ = sentence[i + 1]  
        features.update({  
            '+1:word.lower()': next_word.lower(),  
            '+1:word.istitle()': next_word.istitle(),  
            '+1:postag': next_pos,  
        })  
    else:  
        features['EOS'] = True # End of sentence  
  
    return features
```

```
def sentence_to_features(sentence):  
    return [word_features(sentence, i) for i in range(len(sentence))]
```

```

def sentence_to_labels(sentence):
    return [label for (_, _, label) in sentence]

# Step 3: Prepare training features and labels
X_train = [sentence_to_features(s) for s in train_sentences]
y_train = [sentence_to_labels(s) for s in train_sentences]

# Step 4: Train the Conditional Random Field model
crf_model = sklearn_crfsuite.CRF(
    algorithm='lbfgs',
    c1=0.1,
    c2=0.1,
    max_iterations=100,
    all_possible_transitions=True
)
crf_model.fit(X_train, y_train)

# Step 5: Predict entity labels for a new (unseen) sentence
new_sentence = [("Hello", "O", "O"), ("this", "PRON", "O"), ("is", "AUX", "O"),
                ("Michael", "PROPN", "O"), ("Brown", "PROPN", "O"),
                ("about", "ADP", "O"), ("Order555", "NOUN", "O"),
                ("for", "ADP", "O"), ("MacBookPro", "NOUN", "O")]

X_new = [sentence_to_features(new_sentence)]
predicted_labels = crf_model.predict(X_new)[0]

# Step 6: Display the predicted labels
print("Conditional Random Field: Named Entity Recognition")
print("-----")
print(f'{"Word":<15} {"Predicted Label"}')
print("-" * 30)
for (word, _, _), label in zip(new_sentence, predicted_labels):
    print(f'{"word":<15} {"label"}')

```

Sample Output



Conditional Random Field: Named Entity Recognition

Word Predicted Label

Hello	O
this	O
is	O
Michael	B-CUST
Brown	I-CUST
about	O
Order555	B-ORDER
for	O
MacBookPro	B-PROD

Explanation of Output

Given the new sentence, the trained CRF model correctly tags 'Michael Brown' as a Customer (B-CUST, I-CUST), 'Order555' as an Order ID (B-ORDER), and 'MacBookPro' as a Product (B-PROD), while all other words are labelled 'O' (outside any entity). This demonstrates that the CRF has learned the contextual patterns distinguishing each entity type.

Question 5: Bayesian Network Learning from Data

Problem Statement

An insurance company wants to learn the probabilistic relationships among Age, Income, Vehicle Type, and Insurance Claim directly from historical customer data, rather than specifying the network structure manually, and then use the learned model to predict claim probability for a new customer.

Approach

- Load the historical customer dataset into a Pandas DataFrame (Age, Income, VehicleType, InsuranceClaim).
- Use pgmpy's HillClimbSearch with the BIC (Bayesian Information Criterion) score to learn the best-fitting network structure directly from the data.
- Build a BayesianNetwork from the learned edges and estimate its CPTs using Maximum Likelihood Estimation.
- Run Variable Elimination inference to predict $P(\text{InsuranceClaim} \mid \text{Age, Income, VehicleType})$ for a new customer profile.
- Print the learned graph structure and the final prediction results.

Python Code

```
# Question 5: Bayesian Network Learning from Data
# Learn probabilistic relationships among Age, Income, VehicleType, InsuranceClaim
# Library required: pip install pgmpy pandas

import pandas as pd
from pgmpy.estimators import HillClimbSearch, BicScore, MaximumLikelihoodEstimator
from pgmpy.models import BayesianNetwork
from pgmpy.inference import VariableElimination

# Step 1: Load / create the historical customer dataset using Pandas
# (In practice this would be: data = pd.read_csv("insurance_data.csv"))
data = pd.DataFrame({
    'Age':    ['Young', 'Young', 'Middle', 'Old', 'Old', 'Middle', 'Young', 'Old',
              'Middle', 'Young', 'Old', 'Middle', 'Young', 'Old', 'Middle', 'Young'],
    'Income': ['Low', 'Low', 'High', 'High', 'Low', 'High', 'Low', 'High',
              'Low', 'High', 'Low', 'High', 'Low', 'Low', 'High', 'High'],
    'VehicleType': ['Sedan', 'SUV', 'SUV', 'Sedan', 'Sedan', 'SUV', 'SUV', 'SUV',
                   'Sedan', 'SUV', 'Sedan', 'SUV', 'Sedan', 'Sedan', 'SUV', 'SUV'],
    'InsuranceClaim': ['No', 'Yes', 'Yes', 'No', 'No', 'Yes', 'No', 'Yes',
                      'No', 'Yes', 'No', 'Yes', 'No', 'No', 'Yes', 'Yes']
})
```

```

print("Historical Customer Dataset")
print("-----")
print(data)

# Step 2: Learn the Bayesian Network structure from data using Hill Climb Search
# with the Bayesian Information Criterion (BIC) score
hc_search = HillClimbSearch(data)
best_structure = hc_search.estimate(scoring_method=BicScore(data))

print("\nLearned Bayesian Network Structure")
print("-----")
print("Edges:", list(best_structure.edges()))

# Step 3: Build the Bayesian Network with the learned structure
learned_model = BayesianNetwork(best_structure.edges())

# Step 4: Learn the CPTs (parameters) from data using Maximum Likelihood Estimation
learned_model.fit(data, estimator=MaximumLikelihoodEstimator)

# Step 5: Perform inference - predict probability of InsuranceClaim for a new customer
# New customer: Age = Young, Income = Low, VehicleType = SUV
inference = VariableElimination(learned_model)
result = inference.query(
    variables=['InsuranceClaim'],
    evidence={'Age': 'Young', 'Income': 'Low', 'VehicleType': 'SUV'}
)

# Step 6: Display the learned network structure and prediction results
print("\nPrediction for New Customer (Age=Young, Income=Low, VehicleType=SUV)")
print("-----")
print(result)

```

Sample Output

```

Historical Customer Dataset
-----
      Age Income VehicleType InsuranceClaim
0  Young   Low      Sedan         No
1  Young   Low      SUV          Yes
2  Middle High      SUV          Yes
... (16 rows total)

Learned Bayesian Network Structure
-----
Edges: [('Income', 'InsuranceClaim'), ('VehicleType', 'InsuranceClaim'),
        ('Age', 'Income')]

Prediction for New Customer (Age=Young, Income=Low, VehicleType=SUV)
-----
+-----+-----+
| InsuranceClaim | phi(InsuranceClaim) |
+=====+=====+
| InsuranceClaim(No) | 0.1200 |
+-----+-----+
| InsuranceClaim(Yes) | 0.8800 |
+-----+-----+
```

Explanation of Output

Hill-Climb structure search discovers the most probable dependency graph directly from the data (for example, Income and VehicleType influencing InsuranceClaim, with Age influencing Income). For the new customer (Young, Low Income, SUV), the learned model predicts a high probability of an insurance claim, illustrating how structure learning combined with inference produces actionable, data-driven predictions. Note that exact edges/probabilities depend on the dataset used and may vary slightly between runs.